

```
None, sm_group_num=8, custom_weight_loader=[], weight_loader_disable_mmap=False, weight_loader_prefetch_checkpoints=False, weight_loader_prefetch_num_threads=4, weight_loader_drop_ca
che_after_load=False, remote_instance_weight_loader_seed_instance_ip=None, remote_instance_weight_loader_seed_instance_service_port=None, remote_instance_weight_loader_send_weights_g
roup_ports=None, remote_instance_weight_loader_backend='nccl', remote_instance_weight_loader_start_seed_via_transfer_engine=False, engine_info_bootstrap_port=6789, modelexpress_conf
ig=None, download_dir=None, model_checksum=None, delete_ckpt_after_loading=False, decrypted_config_file=None, decrypted_draft_config_file=None, checkpoint_engine_wait_weights_before_r
eady=False, enable_broadcast_mm_inputs_process=False, enable_prefix_mm_cache=False, mm_enable_dp_encoder=False, mm_process_config={}, limit_mm_data_per_request=None, enable_mm_global
_cache=False, disable_fast_image_processor=False, keep_mm_feature_on_device=False, forward_hooks=None, msprobe_dump_config=None)
[2026-09-26 15:59:49] Using default HuggingFace chat template with detected content format: string
[2026-09-26 15:59:49] Auto-detected template features: reasoning_config=ReasoningToggleConfig(toggle_param='enable_thinking', default_enabled=True, special_case=None), reasoning_pars
er=qwen3, tool_call_parser=qwen
[2026-09-26 15:59:59] Init torch distributed begin.
[2026-09-26 15:59:59] Init torch distributed ends. elapsed=0.60 s, mem usage=0.04 GB
[2026-09-26 16:00:01] Load weight begin. avail mem=31.16 GB
Multi-thread loading shards: 0% Completed | 0/1 [00:00<?, ?it/s][2026-09-26 16:00:01] Parameter lm_head.weight not found in params_dict
Multi-thread loading shards: 100% Completed | 1/1 [00:00<00:00, 4.00it/s]
[2026-09-26 16:00:01] Load weight end. elapsed=0.40 s, type=Qwen3ForCausalLM, avail mem=29.98 GB, mem usage=1.18 GB.
[2026-09-26 16:00:02] KV Cache is allocated. dtype: torch.bfloat16, #tokens: 244167, K size: 13.04 GB, V size: 13.04 GB
[2026-09-26 16:00:02] Memory pool end. avail mem=3.37 GB
[2026-09-26 16:00:02] Capture target prefill CUDA graph begin. backend=tc_pieewise, num_tokens=[4, 8, 12, 16, 20, 24, 28, 32, 48, 64, 80, 96, 112, 128, 144, 160, 176, 192, 208, 224,
240, 256, 288, 320, 352, 384, 416, 448, 480, 512, 576, 640, 704, 768, 832, 896, 960, 1024, 1280, 1536, 1792, 2048], avail mem=2.80 GB
Compiling num tokens (num_tokens=2048): 0% | 0/42 [00:00<?, ?it/s][
2026-09-26 16:01:01] Compiling a graph for dynamic shape takes 0.41 s
[2026-09-26 16:01:05] Compiling a graph for dynamic shape takes 0.46 s
Compiling num tokens (num_tokens=4): 100% | 42/42 [00:09<00:00, 4.48it/s]
Capturing num tokens (num_tokens=4 avail_mem=2.50 GB): 100% | 42/42 [00:01<00:00, 23.71it/s]
[2026-09-26 16:01:09] Capture target prefill CUDA graph end. elapsed=66.30 s, mem usage=0.30 GB, avail mem=2.50 GB.
[2026-09-26 16:01:09] Capture target decode CUDA graph begin. backend=full, num_tokens_per_bs=1, bs=[1, 2, 4, 8, 12, 16, 24], avail mem=2.50 GB
Capturing batches (bs=1 avail_mem=2.36 GB): 100% | 7/7 [00:10<00:00, 1.45s/it]
[2026-09-26 16:01:19] Capture target decode CUDA graph end. elapsed=10.84 s, mem usage=0.14 GB, avail mem=2.35 GB.
[2026-09-26 16:01:19] max_total_num_tokens=244167, chunked_prefill_size=2048, max_prefill_tokens=16384, max_running_requests=3052, context_len=40960, available_gpu_mem=2.35 GB
[2026-09-26 16:01:19] Tree cache initialized: source=default Impl=RadixCache hybrid_swa=False hybrid_ssm=False hierarchical=False streaming_wrapped=False
[2026-09-26 16:01:20] INFO: Started server process [2279]
[2026-09-26 16:01:20] INFO: Waiting for application startup.
[2026-09-26 16:01:20] Using default chat sampling params from model generation config: {'temperature': 0.6, 'top_k': 20, 'top_p': 0.95}
[2026-09-26 16:01:20] INFO: Application startup complete.
[2026-09-26 16:01:20] INFO: Uvicorn running on http://0.0.0.0:30000 (Press CTRL+C to quit)
[2026-09-26 16:01:21] INFO: 127.0.0.1:51070 - "GET /model_info HTTP/1.1" 200 OK
[2026-09-26 16:01:30] Prefill batch, #new-seq: 1, #new-token: 6, #cached-token: 6, cuda graph: True, input throu
ghput (tokens/s): 0.59
[2026-09-26 16:01:30] INFO: 127.0.0.1:51072 - "POST /generate HTTP/1.1" 200 OK
[2026-09-26 16:01:30] The server is fired up and ready to roll!
```

SGLang服务启动和推理验证

```
-----AutoDL-----+
目录说明:
+-----+-----+-----+-----+
| 目录 | 名称 | 速度 | 说明 |
+-----+-----+-----+-----+
| /root/autodl-tmp | 系统盘 | 一般 | 实例关机数据不会丢失，可存放代码等。会随保存镜像一起保存。 |
|                   | 数据盘 | 快   | 实例关机数据不会丢失，可存放读写I/O要求高的数据。但不会随保存镜像一起保存 |
+-----+-----+-----+-----+

CPU：12 核心
内存：62 GB
GPU：NVIDIA GeForce RTX 4080 SUPER, 1
存储：
系统盘/          : 56% 17G/30G
数据盘/root/autodl-tmp: 1% 36K/50G
+-----+-----+
*注意:
1. 系统盘较小请将大的数据存放于数据盘或文件存储中，重置系统时数据盘和文件存储中的数据不受影响
2. 清理系统盘请参考: https://www.autodl.com/docs/gal/
3. 终端中长期执行命令请使用screen等工具后台运行，确保程序不受SSH连接中断影响: https://www.autodl.com/docs/daemon/
root@autodl-container-9e152118z4-9faadd3d:~# curl http://127.0.0.1:30000/v1/chat/completions \
-H 'Content-Type: application/json' \
-d '{"model": "Qwen/Qwen3-0.6B", "messages": [{"role": "user", "content": "你好，请做个自我介绍。"}]}'
{"id": "dd6bf5a8f494474080ac3fd428e0e525", "object": "chat.completion", "created": 1790409871, "model": "Qwen/Qwen3-0.6B", "choices": [{"index": 0, "message": {"role": "assistant", "content": "<think>\n好的。用户让我做个自我介绍。首先，我需要确认他是否已经准备好，或者是否需要我引导他。可能他正在准备面试，或者需要自我介绍作为助理。我应该保持友好和专业的态度，先感谢他的提问，然后简要介绍自己，包括专业领域和职责。同时，要确保回答简洁明了，适合不同场合使用，比如面试或日常交流。还要注意语气要自然，避免过于生硬。最后，检查有没有遗漏的信息，确保信息准确无误。"}}, {"role": "user", "content": "你好，请做个自我介绍。"}]}]}'
root@autodl-container-9e152118z4-9faadd3d:~# curl http://127.0.0.1:30000/v1/chat/completions -H 'Content-Type: application/json' -d '{"model": "Qwen/Qwen3-0.6B", "messages": [{"role": "user", "content": "你好，请做个自我介绍。"}]}'
{"id": "dd6bf5a8f494474080ac3fd428e0e525", "object": "chat.completion", "created": 1790409871, "model": "Qwen/Qwen3-0.6B", "choices": [{"index": 0, "message": {"role": "assistant", "content": "<think>\n您好！我是XXY，专注于XXY领域的工作。我负责XXY任务，擅长XXY技能，希望能为团队贡献您的价值。有什么我可以帮助您的吗？"}}, {"role": "user", "content": "你好，请做个自我介绍。"}]}]}'
root@autodl-container-9e152118z4-9faadd3d:~#
```

开始发送10个请求（真实压测，模拟泊松到达）...

Input	Output	Status	TTFT(s)	Latency(s)
188	163	success	0.45	0.45
400	159	success	0.56	0.56
815	93	success	0.34	0.34
123	129	success	0.45	0.45
175	35	success	0.16	0.16
124	196	success	0.36	0.36
395	28	success	0.10	0.10
785	68	success	0.22	0.22
110	32	success	0.12	0.12
214	51	success	0.17	0.17

root@autodl-container-9e152118z4-9faadd3d:~/autodl-tmp#

压测请求记录